Sthira Demo Script: What to Say, Click and Show

Team ECR: Ojas Mutreja, Vidit Dang, Danica Gupta

Hacktoberfest Hack Day Bengaluru 2026 · Track 2 · 02

2026-10-18

Part 1 The plan on one page

You have two demos to give. The first is two minutes, for the first-round evaluation (14:00 to 15:30). The second is five minutes, for the later rounds (15:45 to 17:30), where the criteria are not published.

The first round is judged on three things. Say each one out loud, in this order.

Table 1. The three first-round criteria, and where you hit each

Criterion	Where you hit it
Clear problem and solution	The first 15 seconds, before you touch the screen
At least one working feature	The live check, then the desktop alert
Relevance to the theme	The learning beat and the close: “Track 2, O2, memory and self-correction, on an open-weight model”

HEADLINE Your AI forgets what you decided. Sthira doesn’t.

The two-minute run, in short:

1. **The problem** (15 s). Agents undo decisions because nothing remembers them.
2. **The wall** (10 s). Show the decisions, drag Rewind back and forward once.
3. **The moment** (30 s). Run the MongoDB check. Let the OVERRULED stamp land.
4. **While you’re away** (25 s). Click “Send an agent to work on the app”. The alert arrives by itself.
5. **It learns** (25 s). Reject one item with a reason. Point at the rule it learned.
6. **The close** (15 s). Open source, local-first, open-weight. Nothing deleted, everything cited.

One rule for both runs: one person talks and one person clicks, so the talker never looks down.

Part 2 Before the judges arrive

Allow ten minutes. Do these in order.

Table 2. Setup checklist

Do	How
Warm the local model	Run `ollama run qwen3:4b "hi"` once, then exit
Open the live site in a tab	https://sthira-mocha.vercel.app (check the model badge shows a model name)
Start the desktop demo (hosted model, fast)	`powershell -ExecutionPolicy Bypass -File scripts\launch-demo.ps1`
Same, fully local (slower, the privacy proof)	Add `-Local` to that command
Settle on a fallback	See Part 5

The launcher wipes and recreates the demo ledger every time, so you always start from the same eight settled decisions. Do not relaunch it in the middle of a session.

The desktop app, the redesign and the reading animation are in pull requests #22, #24, #27 and #28. The commands above need those merged, or the branch `feat/ui-reading-progress` checked out.

Part 3 The two-minute run

0:00 The problem (15 seconds, no screen yet)

Say this. Do not click anything.

“We all make decisions with AI now: ChatGPT, Claude, coding agents. Weeks later an agent cheerfully undoes one, because nothing remembers what was settled, or why. Platform memories are siloed, and none of them track decisions.”

0:15 What it is (10 seconds, show the wall)

“Sthira reads your AI history with an open-weight model, keeps every decision with its reason and your exact words, and never deletes anything.”

Drag the Rewind slider back and forward once, and say: “This is what I had decided on 4 March, and now.” On the demo ledger, 4 March shows one decision, 1 April shows three, and 26 June shows all eight.

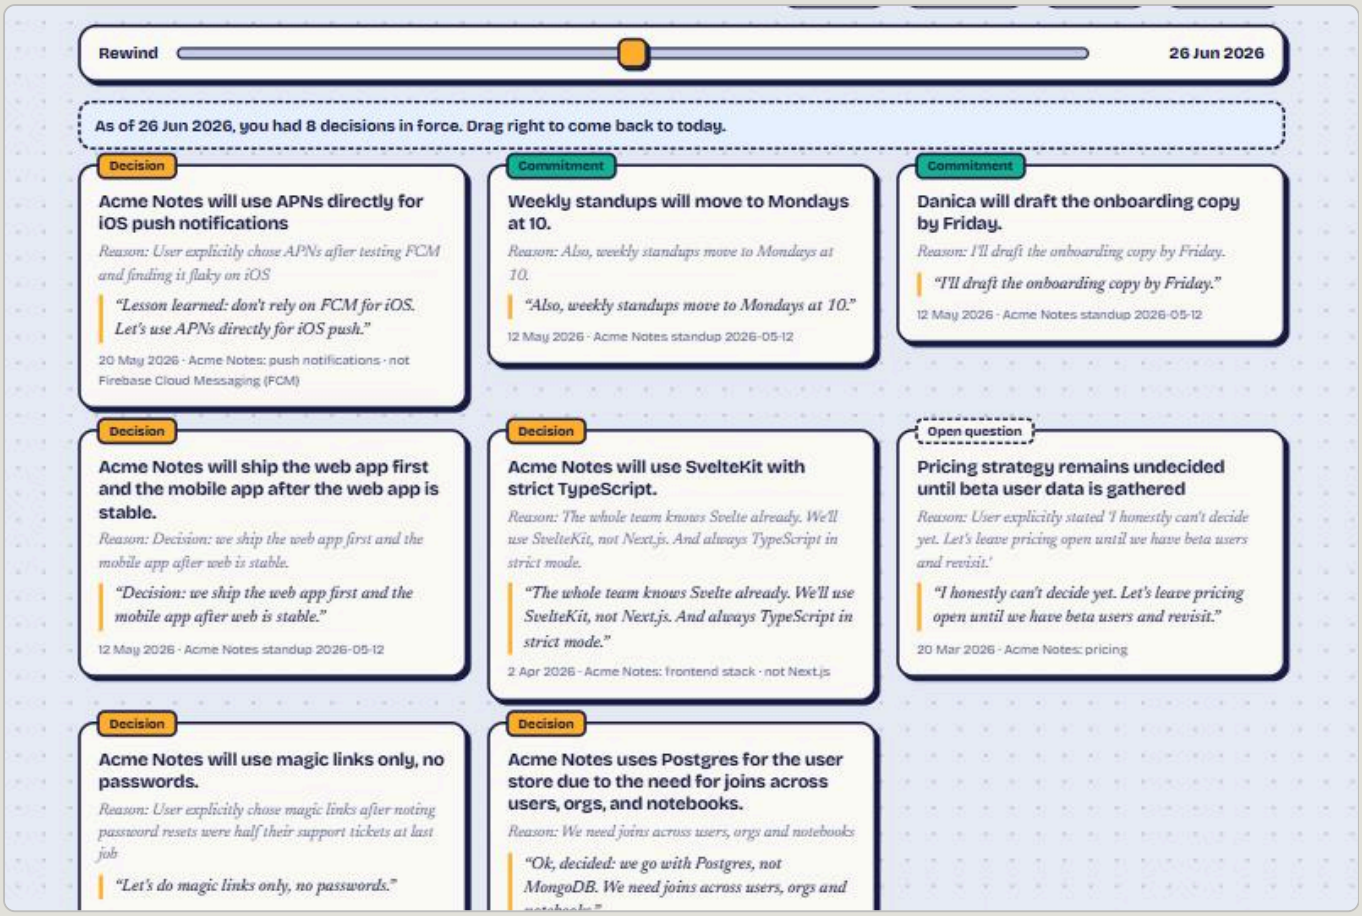


Figure 1. What the wall looks like with Rewind dragged back to 26 June. Each stone shows the decision, the reason, and the exact words you said.

0:25 The moment (30 seconds)

Click the MongoDB example. While the three stage pills light up and the plan types out, narrate: “Search, judge, double-check.” Then stop talking and let the **OVERRULED** stamp land.

“Your agent wanted to migrate to MongoDB. Sthira stops it, shows the exact sentence from March where you ruled that out, and asks if you changed your mind.”

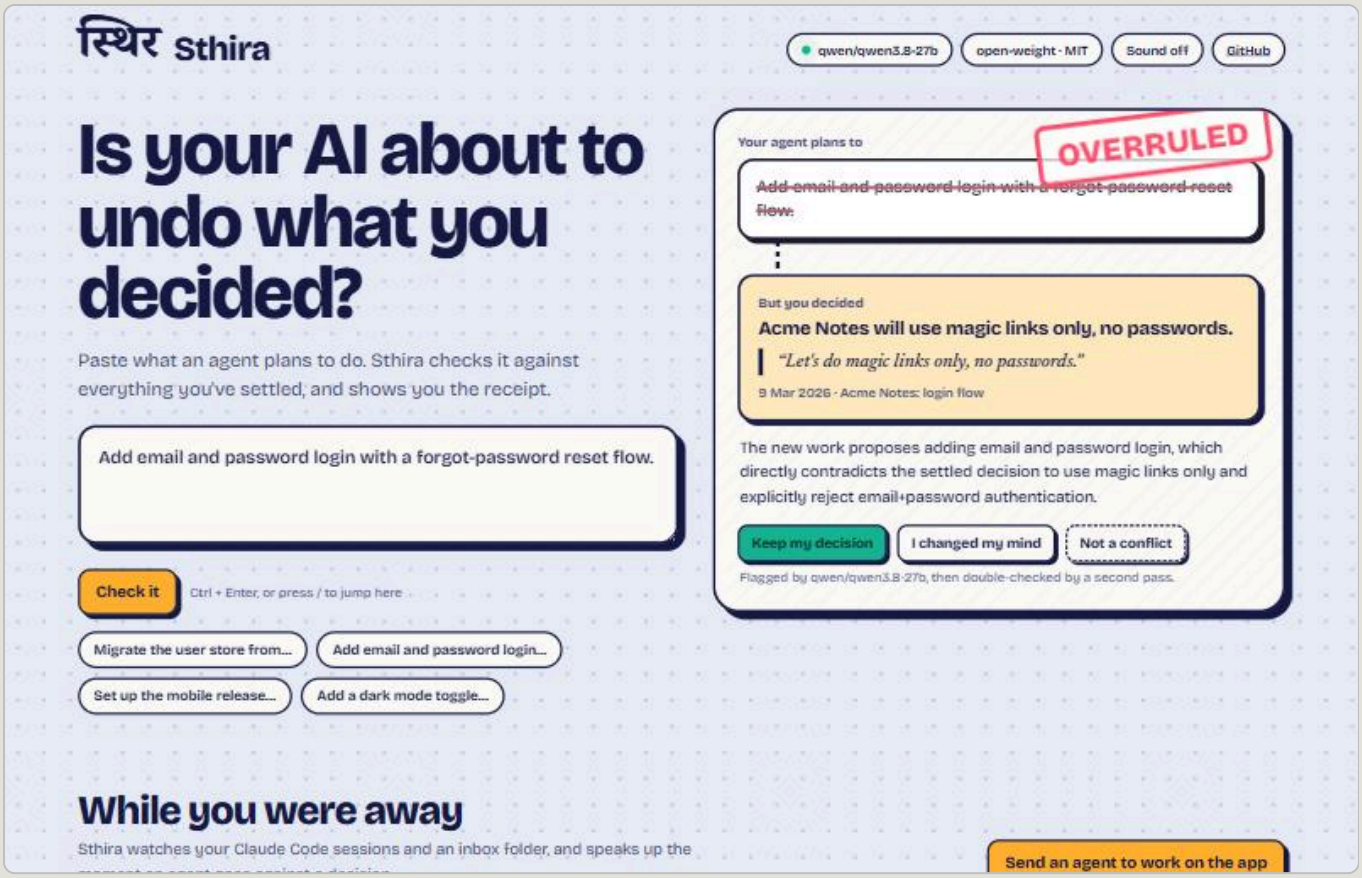


Figure 2. The moment you are building up to. The agent's plan is struck through, and underneath it is the decision you carved in March, with your own words and the date.

0:55 It works while you are away (25 seconds, desktop window)

Click **Send an agent to work on the app**. Keep talking for about ten seconds. Then a red card and a toast arrive on their own.

"Nobody asked it anything. It is watching my Claude Code sessions, caught the agent going against two decisions, and cited both."

1:20 It learns (25 seconds)

Open **To review**. Reject one item and give a reason, for example "small talk is not a decision". Then point at **It learns from you**.

"That reason is now a rule on every future run. On a 4B model, precision on unseen conversations went from 60 percent to 71 percent. And a second critic pass took contradiction checks from 4 of 6 to 6 of 6."

1:45 The close (15 seconds)

“Open source, MIT, local-first, open-weight Qwen3. Nothing is deleted, everything cites its source. Your AI forgets what you decided. Sthira doesn’t.”

Part 4 The five-minute run

Do everything in Part 3, but stretch these five beats.

The reader (60 seconds)

Drag `demo\chatgpt\export-march.json` into **Feed it your history**. Point out the document lighting up line by line, the “2 of 5”, and the estimate of time left.

“That is real progress, not a fake bar. The estimate starts as a guess and corrects itself from the measured time of each chunk.”

Open **To review** and settle one item. It lands with a thunk. Show that its quote is checked against the source file.

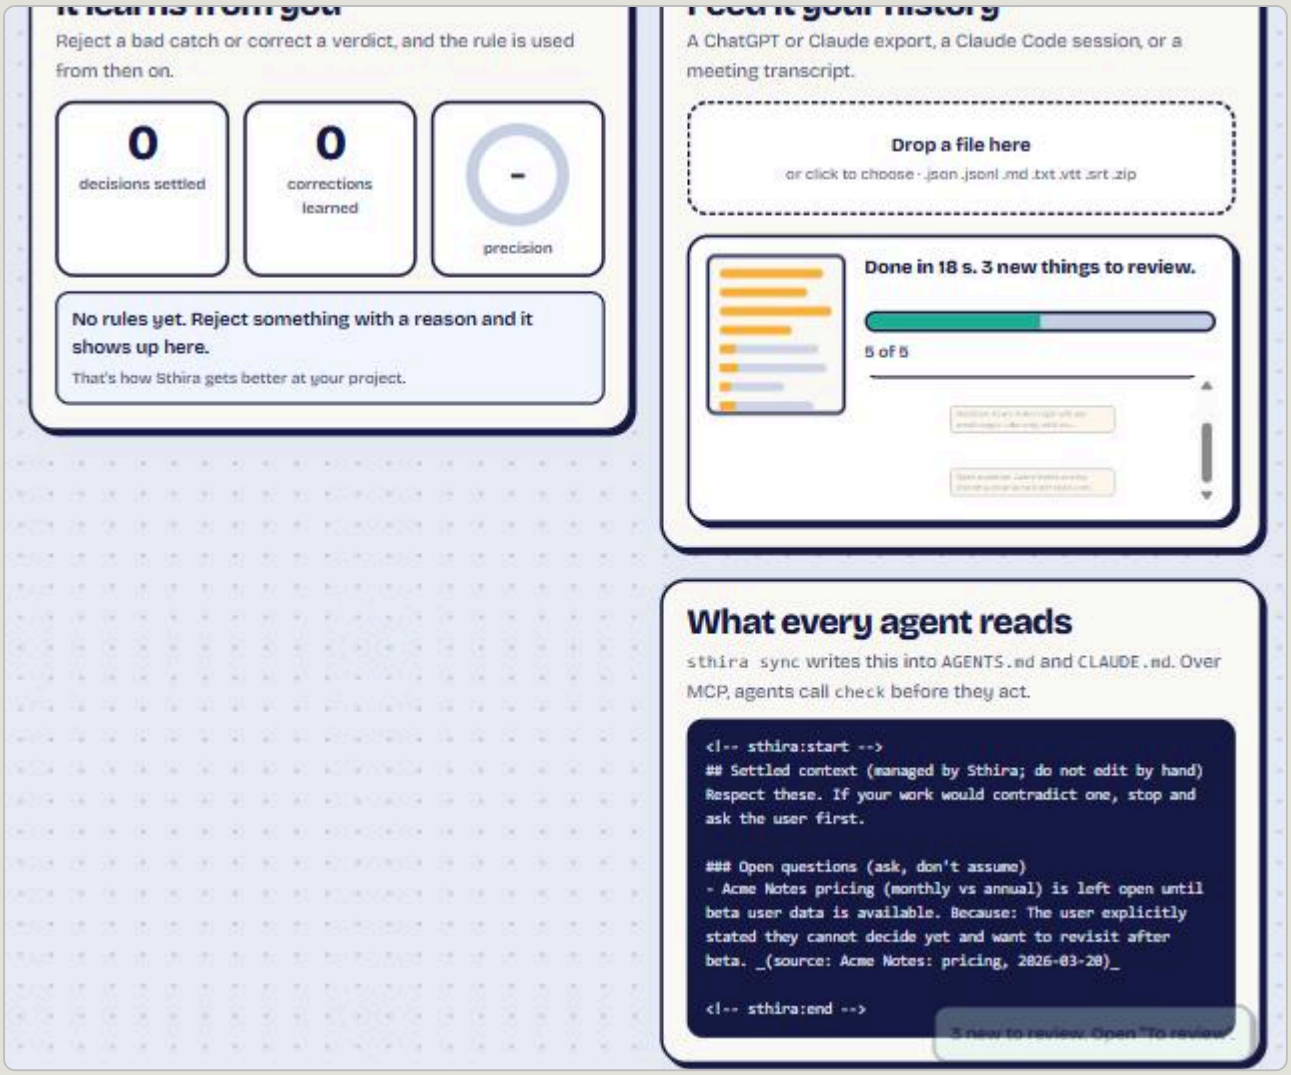


Figure 3. The reader at the end of a run. On a hosted model this takes about two seconds, so run it on the local model if you want the animation to be seen.

I changed my mind (30 seconds)

Run the MongoDB check again and click **I changed my mind**. Give a reason. The old stone cracks away into **History**. Drag Rewind back: “The old decision is still there on 4 March.”

Inside a real agent (60 seconds, only if it is verified)

Do this beat only if `claude mcp list` shows sthira as connected. In Claude Code, type: “Migrate our user store from Postgres to MongoDB. Check with sthira first.” Show Claude calling `check` and stopping to ask you.

If you have not seen it work, skip the beat and show the `AGENTS.md` block at the bottom of the page instead.

We used it on ourselves (30 seconds)

“We pointed Sthira at the session where we built it: 143 messages, 24 entries. It found real decisions, like the project name and dropping the Hacktoberfest-PR angle. It also got things wrong. It logged chatter as ‘the user decided’. That is why a person reviews every entry. And running it found a real bug: rate-limited chunks were skipped silently. We fixed it.”

Honest limits (30 seconds)

Read the list in Part 6 out loud, in your own words. Judges trust a team that says what it cannot do.

Part 5 If something goes wrong

Table 3. Problems and what to do

What you see	Do this
The check says “The model is busy” (the hosted model allows about 4 to 6 checks a minute)	Click one of the four example chips. They fall back to verdicts recorded from a real local run, and the page says so.
The hosted model is down, or no Wi-Fi	Relaunch the desktop demo with `-Local` . It is slower, about 20 to 40 seconds per step, so narrate: “this is the version where nothing leaves the laptop.”
The desktop window is blank or crashed	Switch to the live site tab. It is the same product, minus the background watcher.
The agent alert has not appeared after 20 seconds	Say “it reads sessions once they stop changing”, click once more, and move on. Do not click it five times, because that spends the rate limit.
The reader finishes in two seconds	Say so: “on a laptop CPU this takes about 25 seconds per chunk, which is why the estimate matters.”

What you see	Do this
A judge wants to try their own text	Let them. If the model is offline, a keyword fallback still catches rejected alternatives and labels itself.

Part 6 Numbers you can say, and things you cannot

All numbers below come from one small scenario that we wrote ourselves, run on a laptop. They show the mechanism works. They are not a benchmark.

Table 4. Numbers that are safe to say

Claim	Number
Extraction precision on unseen conversations, qwen3:4b, before and after one round of corrections	0.60 to 0.71 (recall 0.80)
The 8B model on the same scenario	1.0, with nothing to correct
Contradiction checks with the critic pass added (3 real conflicts, 3 traps)	4 of 6 to 6 of 6
Time for the desktop watcher to raise an alert, hosted model	about 9 seconds
Our own build session, read by Sthira	143 messages, 24 entries, with errors we listed
Tests	55 across the open branches; CI on Python 3.11 and 3.13

> FINDING Learning from corrections matters most on small models. That is exactly where it should matter, because small models are what make running privately on a laptop possible.

Do not claim any of the following.

1. That Sthira has been tested inside a live Claude Code session, until it has. So far it has been tested with a real MCP client over the real protocol, and the hook has only been simulated.
2. Any position on the ShipWright leaderboard.
3. That the numbers are a benchmark. They are one small scenario.
4. That Sthira replaces vendor memory. It sits across vendors.
5. That the Gemini importer works on a real export. It is written from the documented format and has not seen one.

Part 7 Questions judges will ask

Why not just use ChatGPT or Claude memory? They are siloed per vendor, other tools cannot read them, and none of them treats decisions, reasons or contradictions as first-class. Sthira works across tools, and the data is yours.

How do you stop made-up memories? Every quote must exist word for word in the source, otherwise the entry is dropped. New entries stay quarantined until a person settles them.

What about memory poisoning? The same answer: nothing becomes settled without a person, and every entry keeps its source.

How is this different from Mem0 and similar tools? They store facts about you. Sthira stores decisions with reasons and history, and checks new work against them. We read about 80 projects before building. The full report is in `docs/research/`.

Why open-weight? Your AI history is private. A 4B model on a laptop is enough to read it, so nothing has to leave your machine. The live site uses a hosted open-weight model, for speed.

How does it meet O2? It retains context across tools. It learns from corrections: rejections and verdict corrections become rules. And it corrects its own mistakes: quotes are verified against the source, and a critic pass re-checks every warning.

What would you build next? Voice notes, Obsidian and git hooks, and warnings when a decision's assumption stops being true. Sharing a ledger across a team comes after that. New sources are small files, so they make easy first contributions, and we have already opened good-first issues.

Part 8 Who does what

Decide this before the first judge arrives, and write it in the table.

Table 5. Roles for the demo

Role	Person
Talks (never looks at the screen)	
Clicks (never talks over the talker)	
Watches the clock and the fallbacks	

And for the question "who built what", agree one honest sentence each. Ojas led the core engine, the desktop app and the interface. Vidit built and reviewed the web and backend work. Danica's part: _____.